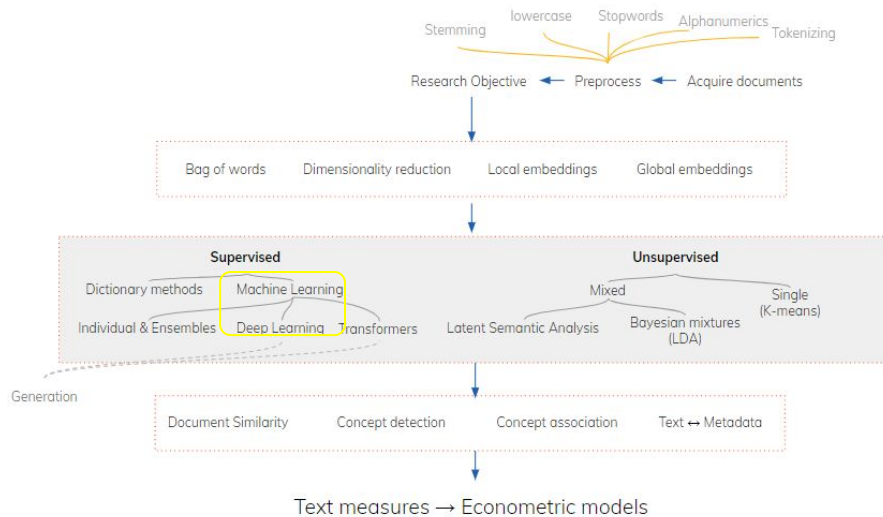


Natural Language Processing in Economics

Advanced Supervised Learning



Neural Networks

Computation Graphs

- Recall the logistic regression model presented in the previous lecture

$$\hat{\mathbf{w}}_{ML} = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{i=1}^N \underbrace{-y_i \log \hat{y}_i - (1 - y_i) \log(1 - \hat{y}_i)}_{\text{Binary Cross Entropy Loss } \mathcal{L}(\hat{y}_i, y_i)}$$

$$\text{with } \hat{y} = f_{\mathbf{w}}(\mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x}) \quad \text{and} \quad \sigma(x) = \frac{1}{1 + e^{-x}}$$

- Minimisation of a **non-linear objective** requires the calculation of gradients $\nabla_{\mathbf{w}} \mathcal{L}(\hat{y}_i, y_i) = (\hat{y}_i - y_i) \mathbf{x}_i$
- This is often a difficult problem to visualise and compute → **computation graphs**

Computation Graphs

Idea

- **Decompose** complex computations into a sequence of atomic assignments, called a **computation graph**
- The **forward pass** takes a training point (\mathbf{x}, \mathbf{y}) as input and computes a loss, ie.

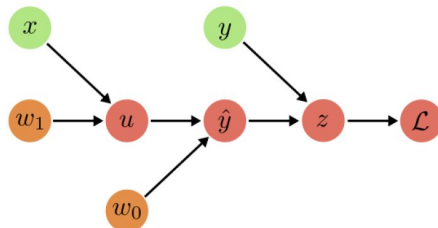
$$\mathcal{L} = -\log p_{model}(y|\mathbf{x}, \mathbf{w})$$

- The relevant gradients can be computed using a **backward pass**, which are efficient due to the use of dynamic programming, ie. storing and reusing intermediate results
- This decomposition and reuse of computation is key to the success of the backpropagation algorithm, the primary workhorse of deep learning

Three kind of nodes

- **Green:** Input nodes
- **Orange:** Parameter nodes
- **Red:** Compute nodes

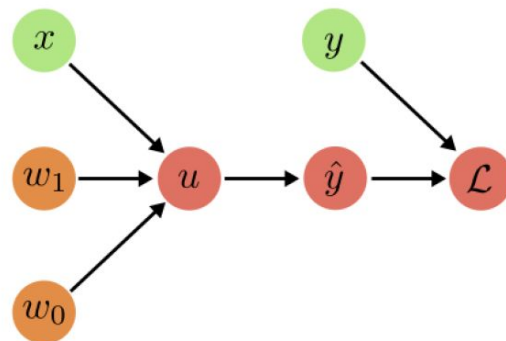
- (1) $u = w_1 x$
- (2) $\hat{y} = w_0 + u$
- (3) $z = \hat{y} - y$
- (4) $\mathcal{L} = z^2$



Computation Graphs

Example - Logistic Regression

- (1) $u = w_0 + w_1x$
- (2) $\hat{y} = \sigma(u)$
- (3) $\mathcal{L} = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$



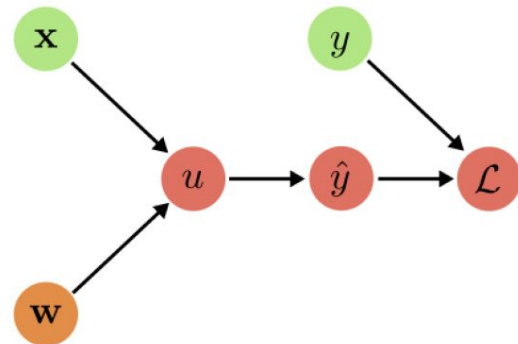
Computation Graphs

Example - Logistic Regression

$$(1) \quad u = \mathbf{w}^\top \mathbf{x}$$

$$(2) \quad \hat{y} = \sigma(u)$$

$$(3) \quad \mathcal{L} = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$




Computation Graphs

Example - Multilayer perceptron

$$(1) \quad \mathbf{h} = \sigma(\mathbf{W}_1^\top \mathbf{x})$$

$$(2) \quad \hat{y} = \sigma(\mathbf{w}_2^\top \mathbf{h})$$

$$(3) \quad \mathcal{L} = -y \log \hat{y} - (1 - y) \log(1 - \hat{y})$$



Computation Graphs

$$\frac{d}{dx}f(g(x)) = \frac{df}{dg} \frac{dg}{dx}$$

$$\frac{d}{dx}f(g_1(x), \dots, g_M(x)) = \sum_{i=1}^M \frac{\partial f}{\partial g_i} \frac{dg_i}{dx}$$

Backpropagation

- We optimise model parameters \mathbf{w} by using gradient descent with respect to the loss function

$$\nabla_{\mathbf{w}} \mathcal{L}(\mathbf{y}, \mathbf{X}, \mathbf{w}) = \nabla_{\mathbf{w}} \sum_{i=1}^N \mathcal{L}(y_i, \mathbf{x}_i, \mathbf{w}) = \sum_{i=1}^N \nabla_{\mathbf{w}} \mathcal{L}(y_i, \mathbf{x}_i, \mathbf{w})$$

Forward Pass:

$$(1) \quad y = x^2$$

$$(2) \quad \mathcal{L} = 2y$$


$$\text{Loss: } \mathcal{L} = 2x^2$$

A simple example of **forward pass** followed by the **backward pass**

Backward Pass:

$$(2) \quad \frac{\partial \mathcal{L}}{\partial y} = \frac{\partial \mathcal{L}}{\partial \mathcal{L}} \frac{\partial \mathcal{L}}{\partial y} = 2$$

$$(1) \quad \frac{\partial \mathcal{L}}{\partial x} = \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial \mathcal{L}}{\partial y} 2x$$



► **Red:** back-propagated gradients

► **Blue:** local gradients

Computation Graphs

$$\frac{d}{dx}f(g(x)) = \frac{df}{dg} \frac{dg}{dx}$$

$$\frac{d}{dx}f(g_1(x), \dots, g_M(x)) = \sum_{i=1}^M \frac{\partial f}{\partial g_i} \frac{dg_i}{dx}$$

Backpropagation

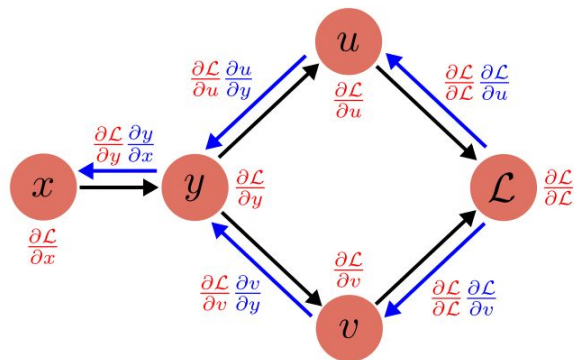
Forward Pass:

- (1) $y = y(x)$
- (2) $u = u(y)$
- (2) $v = v(y)$
- (3) $\mathcal{L} = \mathcal{L}(u, v)$

Backward Pass:

- (3) $\frac{\partial \mathcal{L}}{\partial u} = \frac{\partial \mathcal{L}}{\partial \mathcal{L}} \frac{\partial \mathcal{L}}{\partial u} = \frac{\partial \mathcal{L}}{\partial u}$
- (3) $\frac{\partial \mathcal{L}}{\partial v} = \frac{\partial \mathcal{L}}{\partial \mathcal{L}} \frac{\partial \mathcal{L}}{\partial v} = \frac{\partial \mathcal{L}}{\partial v}$
- (2) $\frac{\partial \mathcal{L}}{\partial y} = \frac{\partial \mathcal{L}}{\partial u} \frac{\partial u}{\partial y} + \frac{\partial \mathcal{L}}{\partial v} \frac{\partial v}{\partial y}$
- (1) $\frac{\partial \mathcal{L}}{\partial x} = \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial x}$

Loss: $\mathcal{L}(u(y(x)), v(y(x)))$



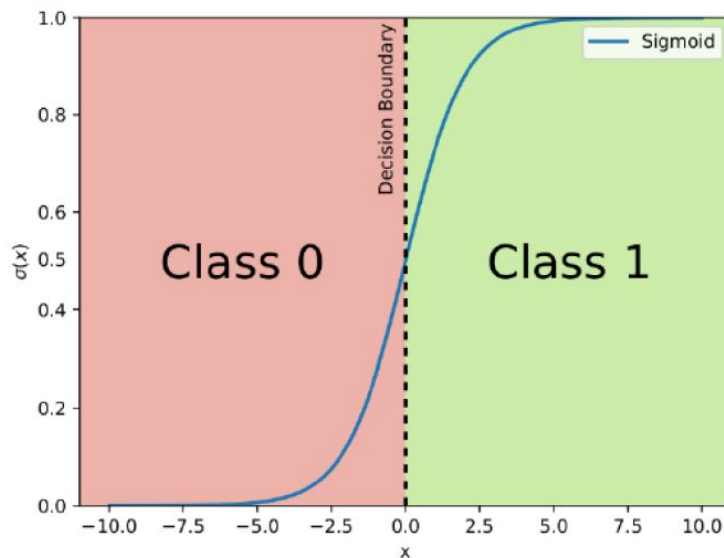
$$\frac{d}{dy}\mathcal{L}(u(y), v(y)) = \frac{\partial \mathcal{L}}{\partial u} \frac{du}{dy} + \frac{\partial \mathcal{L}}{\partial v} \frac{dv}{dy}$$

All incoming gradients must be **summed** up!

A Logistic Regression is a Neural Network

$$\hat{y} = \sigma(\mathbf{w}^T \mathbf{x} + w_0) \quad \text{with } \sigma(x) = \frac{1}{1 + \exp -x}$$

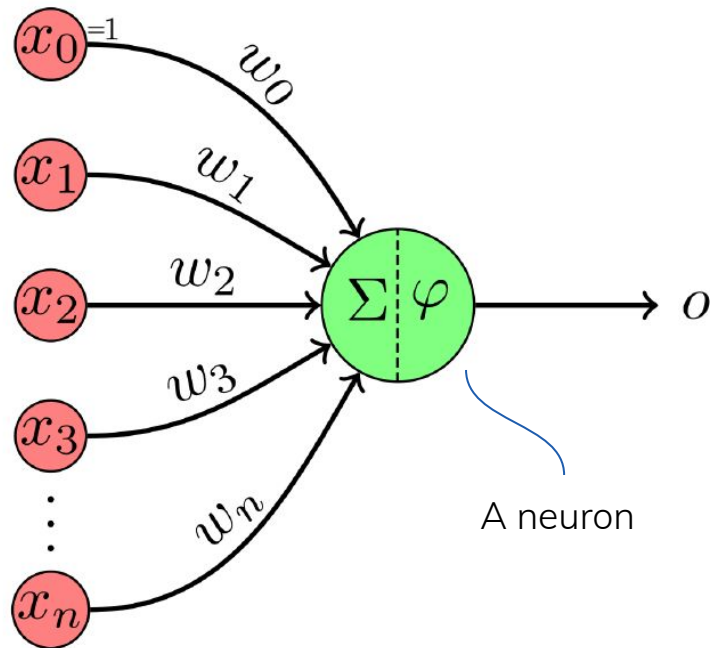
- Let $\mathbf{x} \in \mathbf{R}^2$
- Decision boundary: $\mathbf{w}^T \mathbf{x} + w_0 = 0$
- Decide for Class 0 $\Leftrightarrow \mathbf{w}^T \mathbf{x} < -w_0$
- Decide for Class 1 $\Leftrightarrow \mathbf{w}^T \mathbf{x} > -w_0$



The Perceptron

$$o = \sigma(w_0x_0 + w_1x_1 + w_2x_2 \dots) = \sigma(\sum w_ix_i) = \sigma(\mathbf{w}^T \mathbf{x})$$

- Let $\mathbf{x} \in \mathbb{R}^n$
- Decision boundary: $\sigma(\mathbf{w}^T \mathbf{x}) = 0.5$
- Binary Cross Entropy Loss
- Backpropagation
- Stochastic Gradient Descent



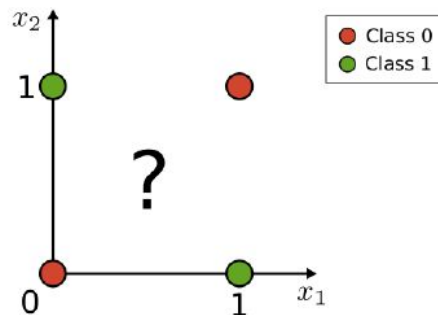
The XOR problem

Linear Classifier:

$$\text{Class 1} \Leftrightarrow \mathbf{w}^T \mathbf{x} > -w_0$$

$$\underbrace{\begin{pmatrix} ? & ? \end{pmatrix}}_{\mathbf{w}^T} \underbrace{\begin{pmatrix} x_1 \\ x_2 \end{pmatrix}}_{\mathbf{x}} > \underbrace{?}_{-w_0}$$

x_1	x_2	$\text{XOR}(x_1, x_2)$
0	0	0
0	1	1
1	0	1
1	1	0

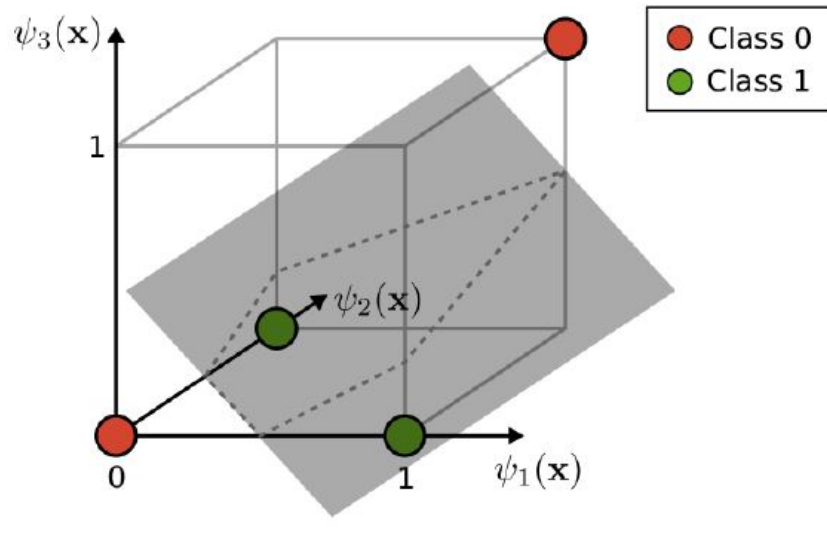


- Minsky & Papert's 1969 book on the limitations of perceptrons
- The first AI winter

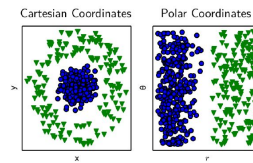
The XOR solution

Linear classifier with
non-linear features ψ :

$$\mathbf{w}^\top \underbrace{\begin{pmatrix} x_1 \\ x_2 \\ x_1 x_2 \end{pmatrix}}_{\psi(\mathbf{x})} > -w_0$$



- Non-linear features allow linear classifier to solve non-linear classification problems
- Analogous to using kernel tricks in regressions or SVMs



Universal Approximation Theorem

Theorem (Cybenko, 1989; Hornik, 1991)

A feedforward neural network with a single hidden layer containing a finite number of neurons can approximate any continuous function on a compact subset of \mathbf{R}^n

- This result is agnostic about the choice of activation function
- The result does not imply the approximation (weights) can be learned

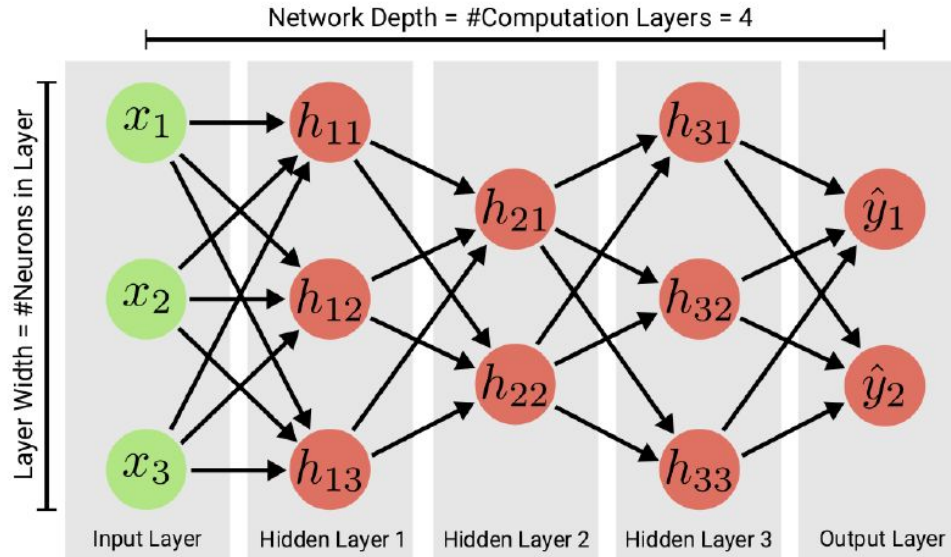
Multilayer Perceptrons

- MLPs are **feedforward** neural networks (no feedback connections)
- They compose several nonlinear functions $\mathbf{f}(\mathbf{x}) = \hat{\mathbf{y}}(\mathbf{h}_3(\mathbf{h}_2(\mathbf{h}_1(\mathbf{x}))))$ where \mathbf{h} are hidden layers and \mathbf{y} is the output layer
- The data specifies only the behavior of the output layer (thus the name **hidden**)
- Each layer \mathbf{i} comprises multiple neurons \mathbf{j} which are implemented as **affine transformations** ($\mathbf{a}^\top \mathbf{x} + \mathbf{b}$) followed by nonlinear activation functions \mathbf{g}

$$h_{ij} = g(\mathbf{a}_{ij}^\top \mathbf{h}_{i-1} + \mathbf{b}_{ij})$$


- Each neuron in each layer is fully connected to all neurons of the previous layer
- The overall length of the chain is the **depth** of the model \rightarrow Deep Learning

Multilayer Perceptrons

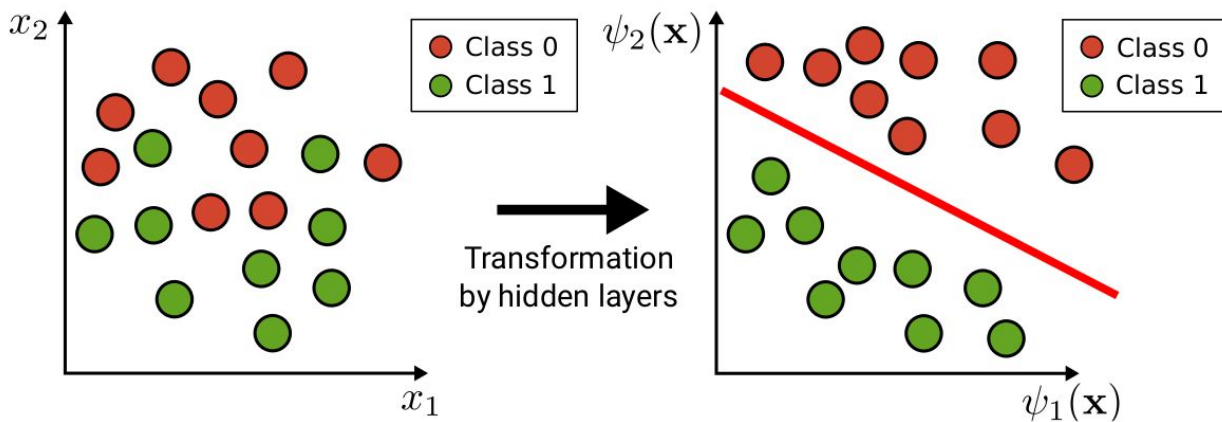


- Neurons are grouped into layers, each neuron is **fully connected** to all previous ones
- **Hidden layer h** has an activation function **g** and weights **A, b**


Multilayer Perceptrons - Feature learning



$$\mathbf{h}_3 = \psi(\mathbf{x})$$



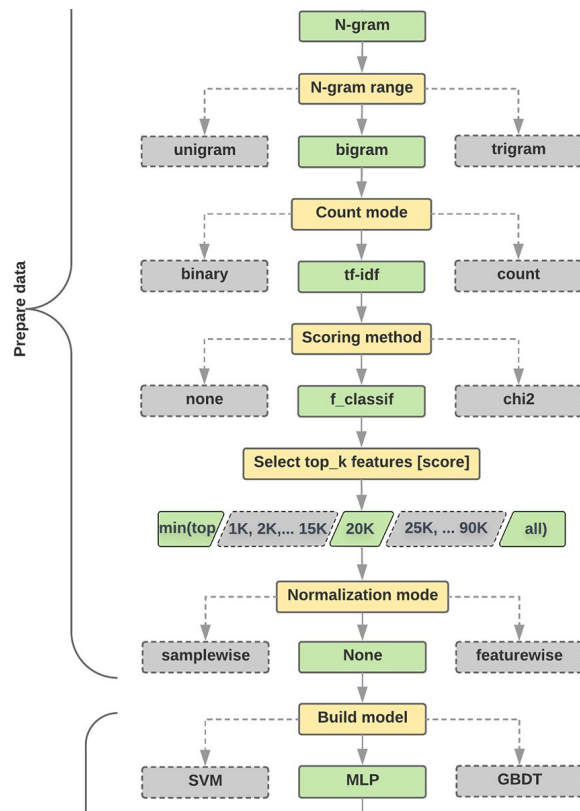
Multilayer Perceptrons

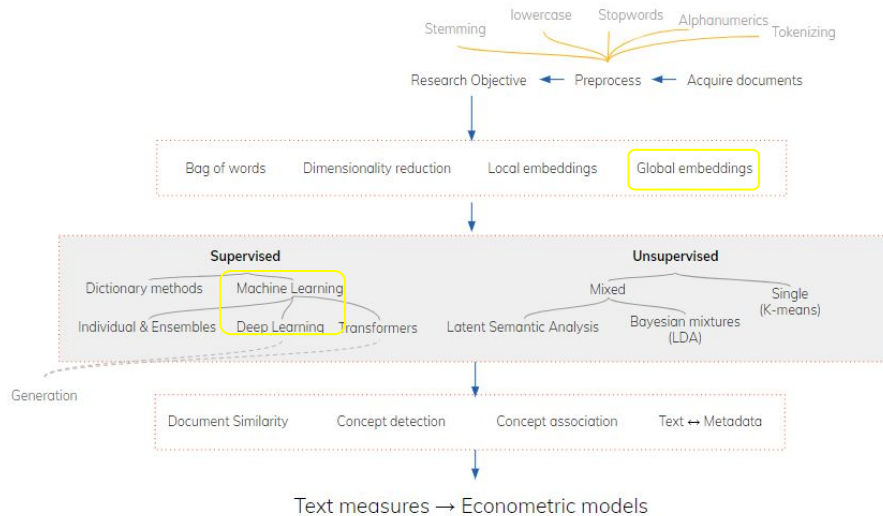


Google Developers Guide for Text Classification

- With relative few, longer documents, use an MLP
- The inputs \mathbf{x} are **tf-idf weighted bigrams**
 - Select 20,000 features using SML
- Employ an MLP with two hidden layers

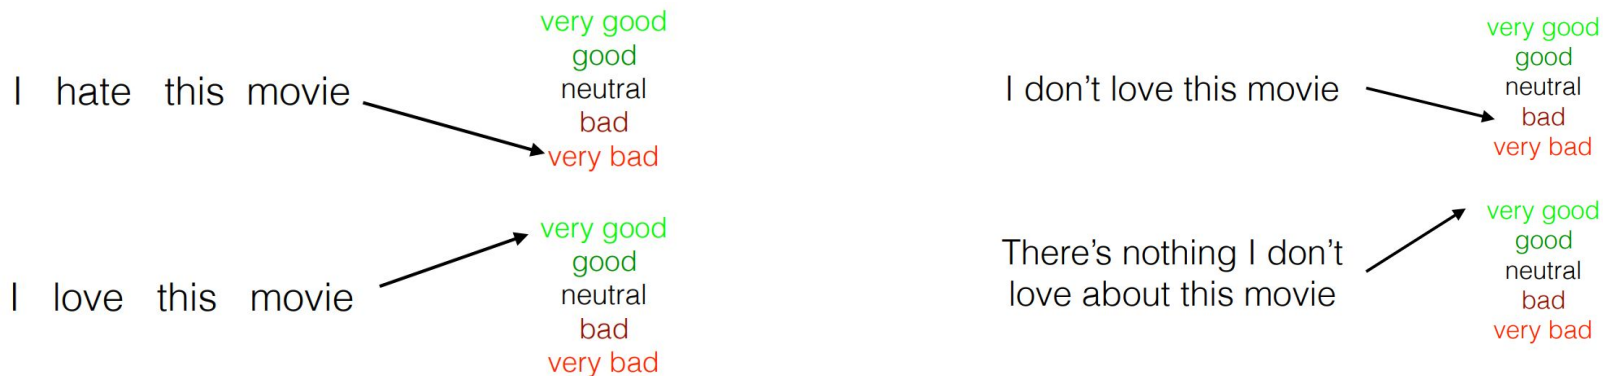
A MLP is one of the models tried by the Peterson and Spirling (2018) paper we saw on predicting polarisation in the UK parliament





Sequential Models

Classical IMDb classification problem



- Bag of words models can't capture the importance of "don't love" or "nothing don't love", even with many hidden layers
- N-grams have a large feature space and don't share information across similar words and n-grams

Not addressed: Deep CBOW, Autoencoders, Convolutional networks

Sequence data

The deep learning breakthrough on NLP → move from bag-of-many representations to sequences

- Rather than inputting counts over n-grams, these models take as input a sequence of tokens

Note as well that feedforward MLP handle fix-length data, but many applications have variable lengths

$$\mathbf{x}_1 = (\mathbf{x}_{1,1}, \mathbf{x}_{1,2}, \mathbf{x}_{1,3}, \mathbf{x}_{1,4}, \mathbf{x}_{1,5})$$

$$\mathbf{x}_2 = (\mathbf{x}_{2,1}, \mathbf{x}_{2,2}, \mathbf{x}_{2,3}, \mathbf{x}_{2,4}, \mathbf{x}_{2,5})$$

$$\mathbf{x}_3 = (\mathbf{x}_{3,1}, \mathbf{x}_{3,2}, \mathbf{x}_{3,3}, \mathbf{x}_{3,4}, \mathbf{x}_{3,5})$$

vs

$$\mathbf{x}_1 = (\mathbf{x}_{1,1}, \mathbf{x}_{1,2}, \mathbf{x}_{1,3}, \mathbf{x}_{1,4}, \mathbf{x}_{1,5})$$

$$\mathbf{x}_2 = (\mathbf{x}_{2,1}, \mathbf{x}_{2,2}, \mathbf{x}_{2,3})$$

$$\mathbf{x}_3 = (\mathbf{x}_{3,1}, \mathbf{x}_{3,2}, \mathbf{x}_{3,3}, \mathbf{x}_{3,4})$$

Sentence sentiment classification

Sound phenomene Recognition

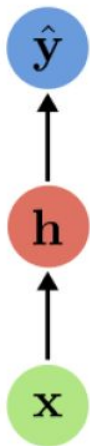
Filmed activities

Sub-par solutions

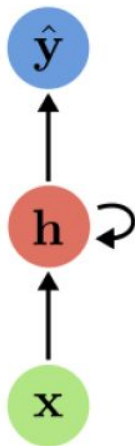
- Zero-pad up to length of longest sequence
- Iterate a single layer over each value

Recurrent Neural Network

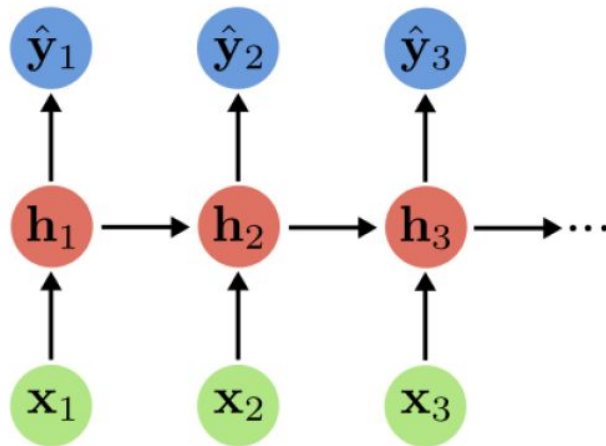
Feedforward
Neural Network



Recurrent Neural Network (RNN)
with feedback connection



Recurrent Neural Network (RNN)
unrolled over time (index = time t)



Core idea

- Update hidden state h based on a new input and the previous hidden state using the same parameters at each time step \rightarrow **encodes** sequences into vectors, & **decodes** vectors into sequences
- Allows for processing **sequences of variable length**, and stores long-term dependencies

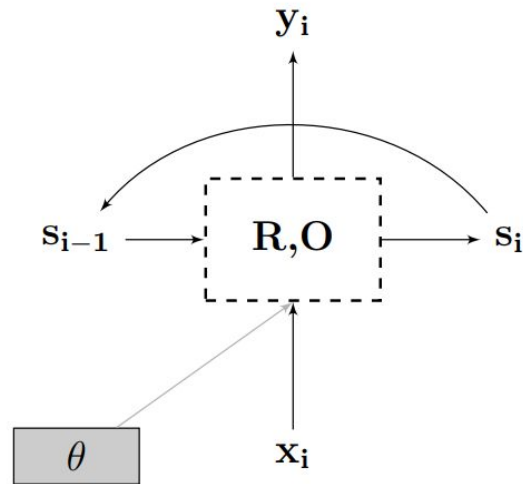
Recurrent Neural Network Architecture

- At each step t :
 - A recursion function $R(s_{t-1}, x_t; \theta_R)$ computes the state vector s_t given current word x_t and previous state s_{t-1}
 - An output function $O(s_t; \theta_O)$ computes the state vector y (to be compared to the outcome variable in the dataset).

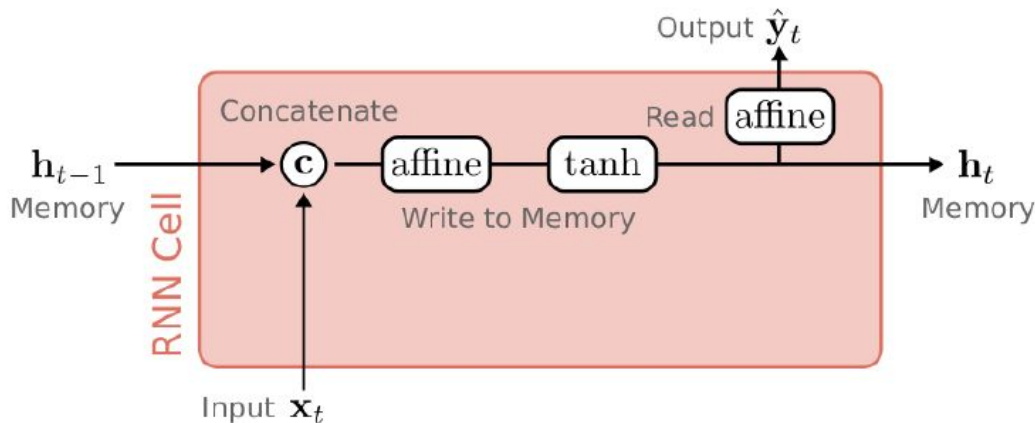
$$\hat{y}_t = O(s_t, \theta_O)$$

$$s_t = R(s_{t-1}, x_t, \theta_R)$$

The parameters of those functions $\theta = (\theta_R, \theta_O)$ are learned during training



Recurrent Neural Network Architecture

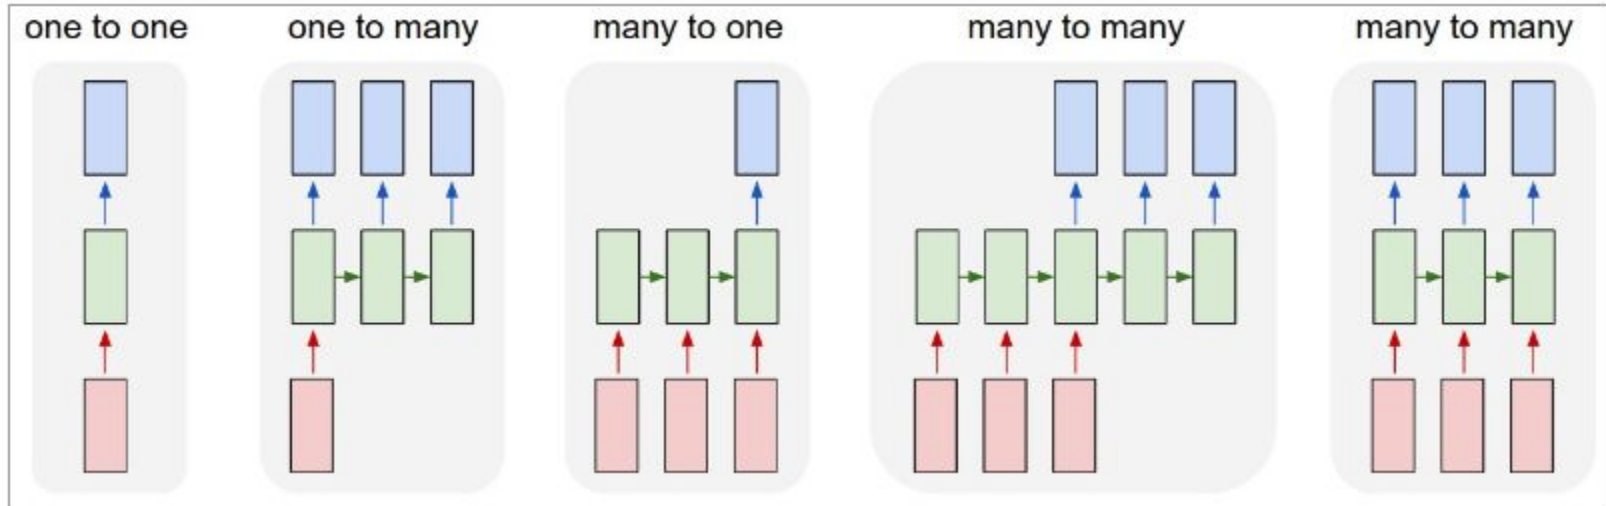


$$\mathbf{h}_t = f_h(\mathbf{h}_{t-1}, \mathbf{x}_t) = \tanh \left(\mathbf{A}^{(h)} \mathbf{h}_{t-1} + \mathbf{A}^{(x)} \mathbf{x}_t + b \right)$$

$$\hat{\mathbf{y}}_t = f_y(\mathbf{h}_t) = \mathbf{A}^{(y)} \mathbf{h}_t$$

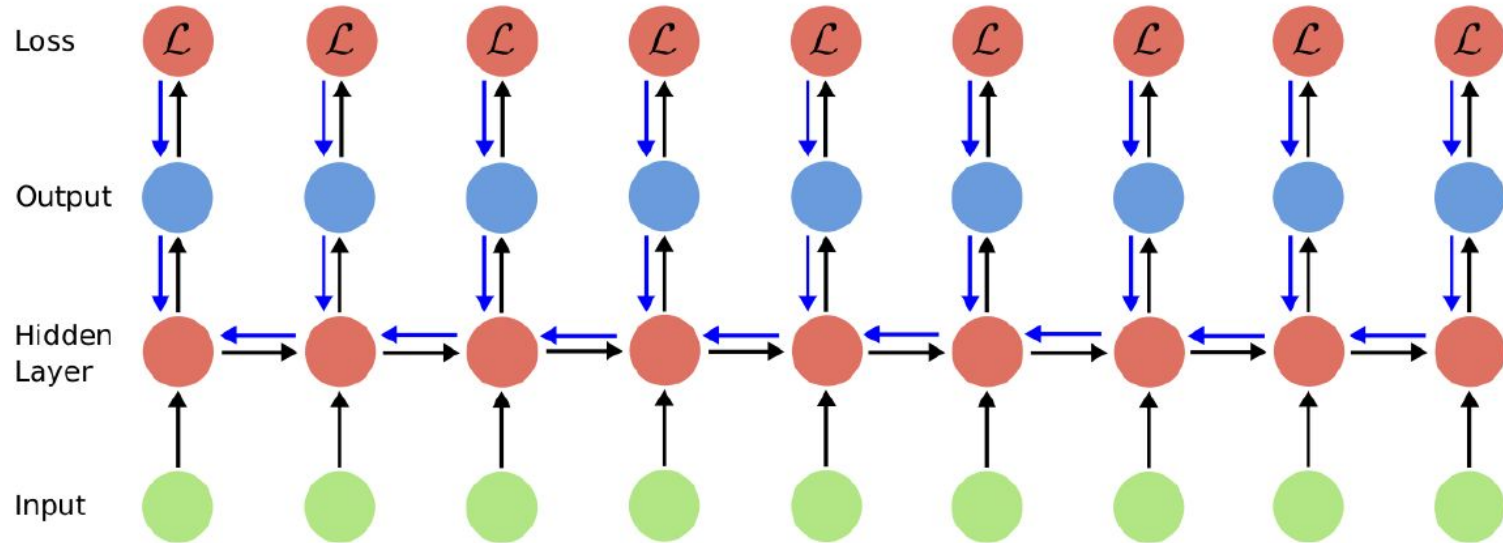
Parameters $\mathbf{A}^{(h)}$, $\mathbf{A}^{(x)}$, $\mathbf{A}^{(y)}$ and b are constant over time.

Recurrent Neural Network Architecture




- **One to one:** image to label
- **One to many:** image to sentence
- **Many to one:** sentence to label
- **Many to many:** sentence to sentence

Training a Recurrent Neural Network



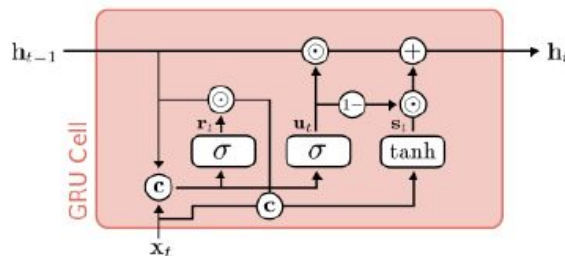
Gated Recurrent Neural Network

UGRNN




Collins, 2017

GRU



Cho, 2014

LSTM



Hochreiter, 1997


Shared attribute: **gates** for filtering information

- **UGRNN**: A sigmoid gate controls how much new information is added
- **GRU**: A reset gate controls the relevance of the previous state on **f**
- **LSTM**: An additional cell **c** carries long term information

Gated Recurrent Neural Network

Caveats


- **Sequence-to-sequence mapping:** Traditional RNNs are designed for synchronized input-output mapping. In many applications, such as machine translation or text summarising, often we deal with variable-length outputs
- **Information bottleneck:** in traditional RNNs, the hidden state must encode all the relevant information from the input sequence into the fixed-length vector, leading to memory leakage and forgetting long-term dependencies
- **Bidirectional context:** RNNs process sequences in a forward-only manner, which means they have limited access to future context when making predictions. This can be problematic for tasks where future context matters



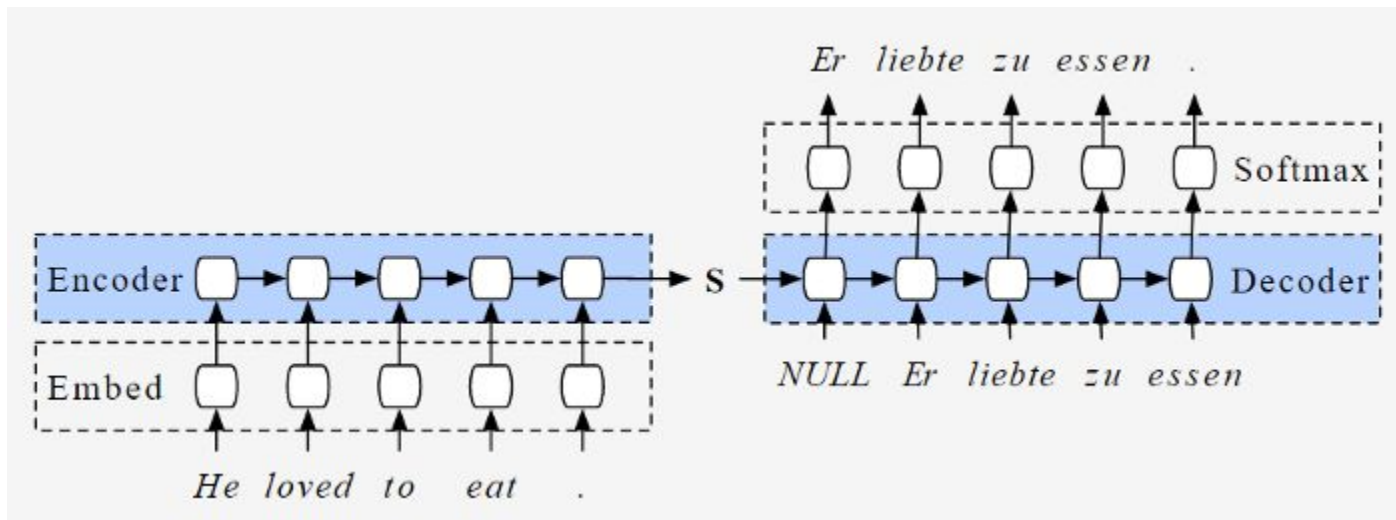
Encoder-Decoder models

Encoder-Decoders: Machine Translation

- The **seq2seq** model:
 - $x^{(i)}$; i^{th} word input
 - $y^{(i)}$; i^{th} word output
 - c ; context vector
- Applications
 - Question Answering
 - Machine translation
 - Machine conversation



Encoder-Decoders: Machine Translation



Problem: The **s** vector aggregates all contextual relations within a paragraph, and does not discriminate among within-string relations

Solution: The **Attention** mechanism

Attention mechanism

- Probabilistic retrieval of a **value** v_i for a **query** q based on a **key** k_i

$$\text{similarity}(q_m, k_i) = q_m^\top k_i$$

- Similarities are transformed into probabilities through a softmax

$$a_{m,i} = \frac{\exp(q_m^\top k_i)}{\sum_j \exp(q_m^\top k_j)}$$

- The output is a linear combination of values or vectors v_i

$$\text{attention}(q_m, \mathbf{k}, \mathbf{v}) = \sum_i a_{m,i} \cdot v_i$$

- The embeddings $\mathbf{q}, \mathbf{k}, \mathbf{v}$ are learned from the data

Attention mechanism in Machine Translation

- **values** v : The **encoder** hidden states h_i for $i \in \{1, 2, \dots, H\}$
- **queries** q : The **decoder** hidden states s_i for $i \in \{1, 2, \dots, S\}$
- **keys** k : The **encoder** hidden states h_i for $i \in \{1, 2, \dots, H\}$

- The similarity measure for decoding step i and encoding step j is

$$\text{similarity}(s_i, h_j) = s_i^\top h_j$$

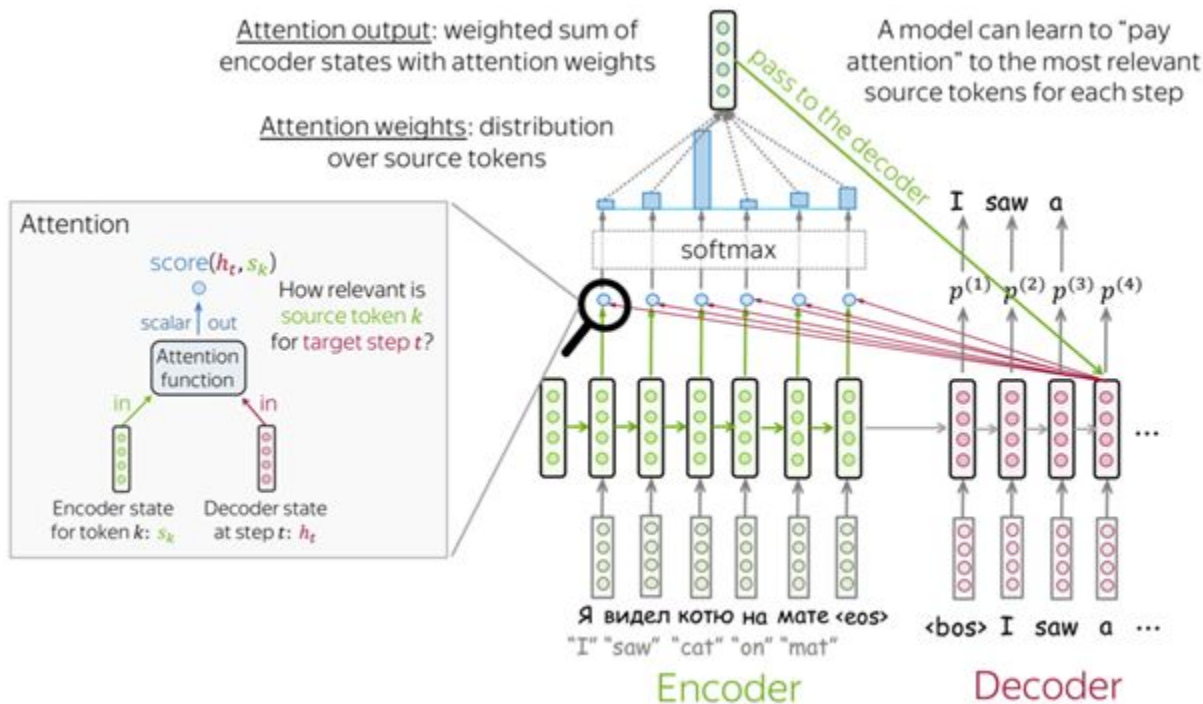
- The attention weight is the softmax probability

$$a_{i,j} = \frac{\exp(s_i^\top h_j)}{\sum_j \exp(s_i^\top h_j)}$$

- The new context vector for decoding step i is the **attention**

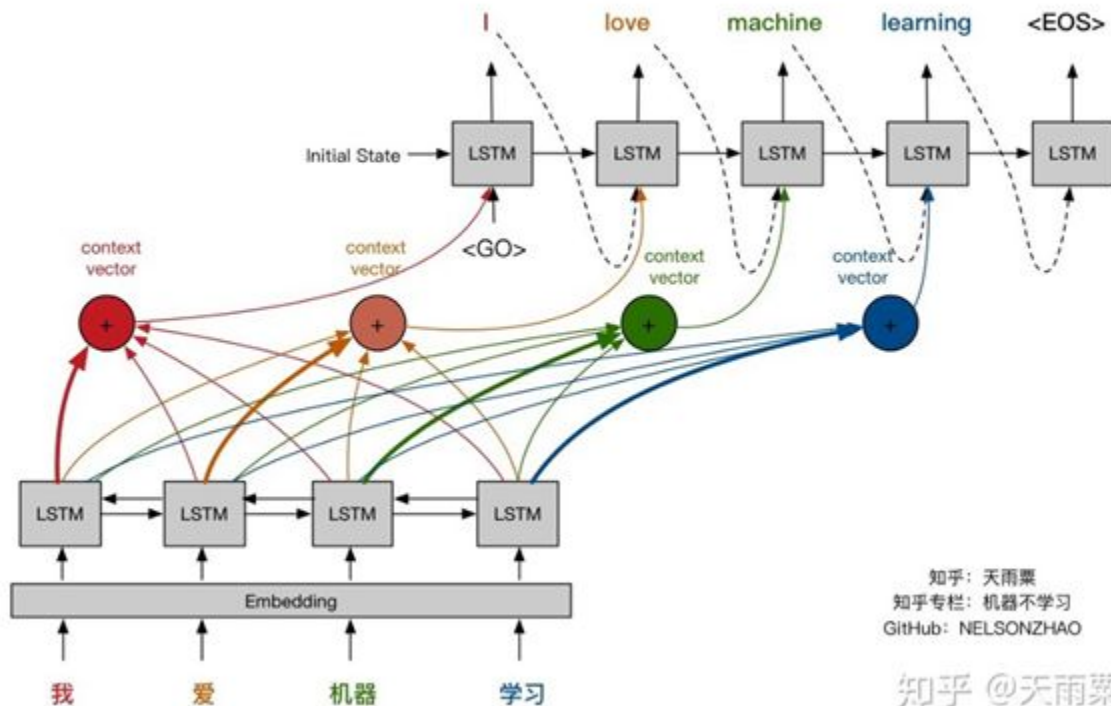
$$c_i = \sum_j a_{i,j} \cdot h_j$$

Attention mechanism in Machine Translation

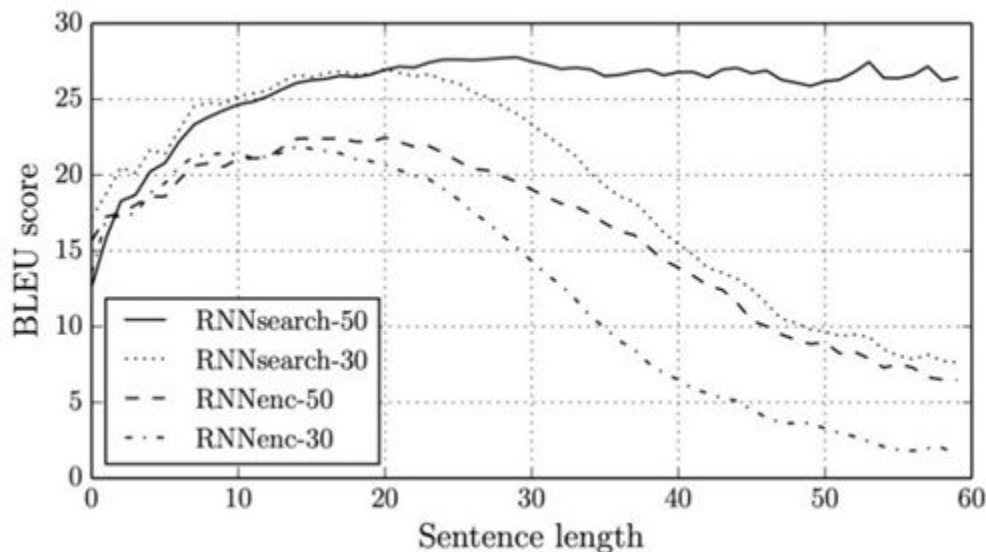


A layer, like a dense layer, an embedding layer


Attention mechanism in Machine Translation

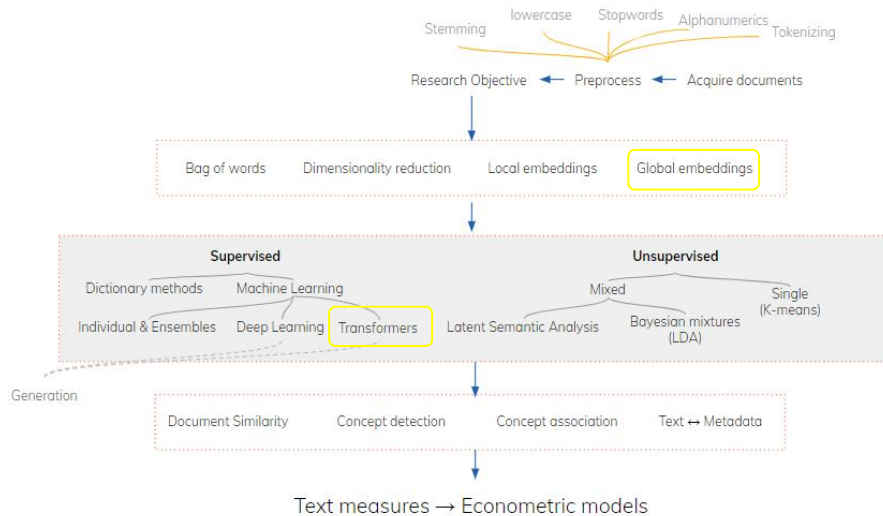


Attention mechanism in Machine Translation



The Bilingual Evaluation Understudy score is a metric used to evaluate the quality of machine-generated translations by comparing them to one or more human reference translations.





Transformers

Transformers

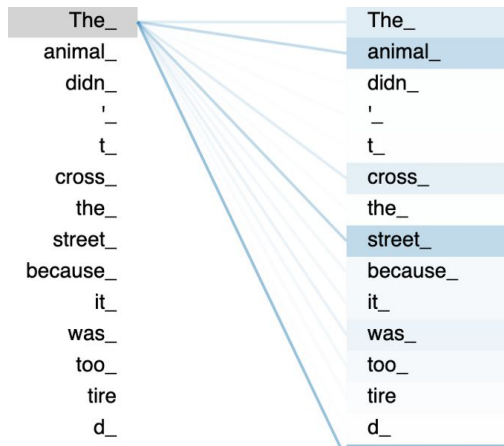

- Most NLP deep learning use transformer-based architectures since [Vaswani et al. \(2017\)](#)
- Recurrent neural nets can process whole documents word-by-word, but they have to sweep through the whole document at each training epoch. This greatly reduces their scalability, as it creates I/O bottlenecks and prevents parallelization.

Transformers overcome this limitation:

- These methods process inputs in parallel, making them highly scalable and efficient for large-scale NLP tasks. They don't process the entire input sequentially (no recurrence).
- They use self-attention mechanisms, enabling them to capture long-range dependencies in text effectively. This allows to retain relationships between words or tokens that are distant from each other.

Attention in AI Research

- **Attention in Computer Vision (2014)**
 - Used to highlight important parts of an image that assist downstream tasks
- **Attention in Natural Language Processing**
 - 2015: Aligned Machine Translation
 - 2017: Transformer Networks & Language Modeling



Transformers in NLP

- Transformer models consist of (many) stacked blocks of parallel **attention heads**
 - These are machine-reading filters, which allow each word to scan over every other word in the document and pick up the most predictive key-value interactions
- **Encoder models**
 - **BERT** - Trained to predict left-out words in the middle of a sequence
 - **AlphaFold** - Protein folding encoder
 - **CLIP** - Image & Object classification
- **Decoder models**
 - **GPT** - Train a transformer to predict the next word at the end of a sequence
 - **LLaMA** - Commonsense reasoning, reading comprehension
 - **Alpaca** - Text generation and classification task

Bottom line: many pre-trained models, can be quickly fine-tuned for downstream tasks ([Transfer Learning](#))

huggingface 🤗

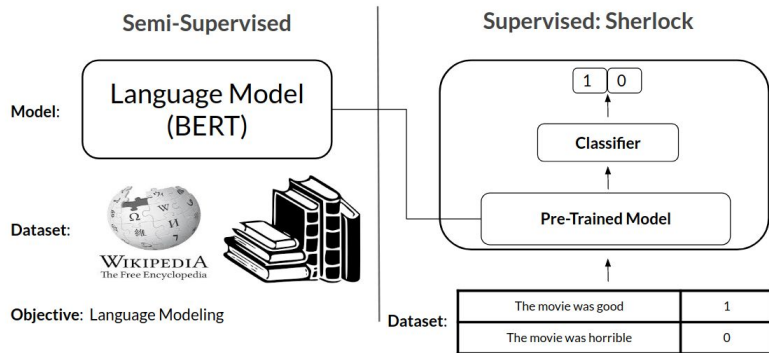
- A popular open-source library and platform for NLP tasks. It offers a comprehensive set of tools and deep learning resources, including Python's most popular **transformer** library.
- It contains hundreds of pre-trained models based on most popular transformer architectures, trained on massive amounts of text data and easily fine-tunable.
- Easy to use API interface, and a large repository of community-contributed pre-trained models, datasets, and evaluation metrics.
- A Hub where people upload and share their models, and explore others'.

```
from transformers import pipeline
sentiment_analysis = pipeline("sentiment-analysis")

pos_text = "I enjoy studying computational algorithms."
neg_text = "I dislike sleeping late everyday."

pos_sent = sentiment_analysis(pos_text)[0]
print(pos_sent['label'], pos_sent['score'])

neg_sent = sentiment_analysis(neg_text)[0]
print(neg_sent['label'], neg_sent['score'])
```



Self-attention mechanism

- A self-referencing attention mechanism
 - **values** v: Word embedding vectors \mathbf{x} (either input or output)
 - **keys** k: Word embedding vectors \mathbf{x} (either input or output)
 - **queries** q: Word embedding \mathbf{x}_i of the previously generated output
- The similarity measure for **queried embedding** \mathbf{x}_i and other embeddings is

$$h_i = \sum_{j=1}^{n_L} a(x_i, x_j) x_j$$

- The attention weight is the softmax probability (sum of weights = 1)

$$a(x_i, x_j) = \text{softmax}\left(\frac{x_i \cdot x_j}{\sqrt{n_E}}\right) = \frac{\exp\left(\frac{x_i \cdot x_j}{\sqrt{n_E}}\right)}{\sum_{k=1}^{n_L} \exp\left(\frac{x_i \cdot x_k}{\sqrt{n_E}}\right)}$$

Scaled by embedding length

- **Note**
 - Basic self-attention has no learnable parameters & ignores word order

Self-attention mechanism - example

- Consider a sentence

the, cat, walks, on, the, street

- with embeddings

$\mathbf{x}_{\text{the}}, \mathbf{x}_{\text{cat}}, \mathbf{x}_{\text{walks}}, \mathbf{x}_{\text{on}}, \mathbf{x}_{\text{the}}, \mathbf{x}_{\text{street}}$

- Feeding this sentence into **self-attention layer** produces

$\mathbf{h}_{\text{the}}, \mathbf{h}_{\text{cat}}, \mathbf{h}_{\text{walks}}, \mathbf{h}_{\text{on}}, \mathbf{h}_{\text{the}}, \mathbf{h}_{\text{street}}$

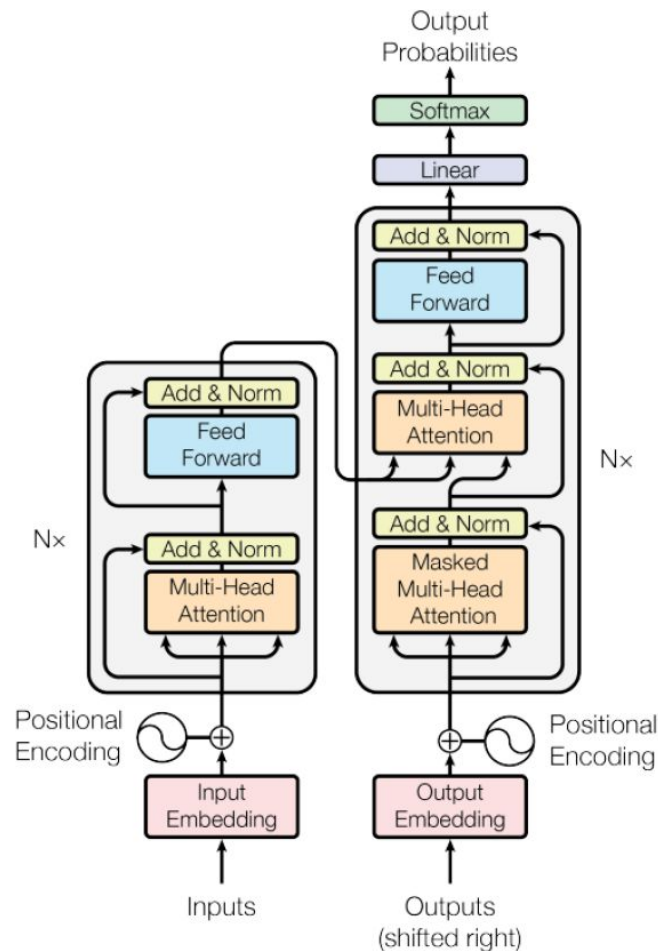
where

$$\mathbf{h}_i = a(x_i \cdot \mathbf{x}_{\text{the}})\mathbf{x}_{\text{the}} + a(x_i \cdot \mathbf{x}_{\text{cat}})\mathbf{x}_{\text{cat}} + \dots + a(x_i \cdot \mathbf{x}_{\text{street}})\mathbf{x}_{\text{street}}$$

Paper

Vaswani et al. (2017) - Attention Is All You Need

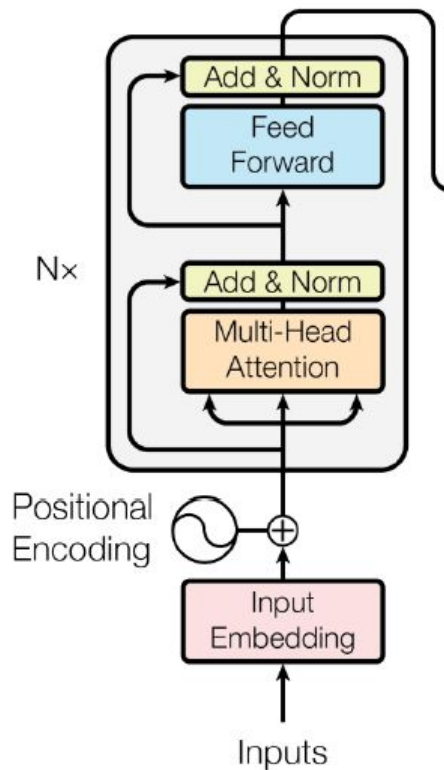
- An **encoder-decoder** based on attention
- No recurrence
 - **Input:** entire sequence to translate
 - **Output:** entire translated sentence
 - **Positional encoding** necessary to distinguish embeddings from some BoW representation
- **Core idea:** use attention to learn abstract embeddings
 - **n=1:** pairs of embeddings
 - **n=2:** pairs of pairs of embeddings
 - ...



Paper

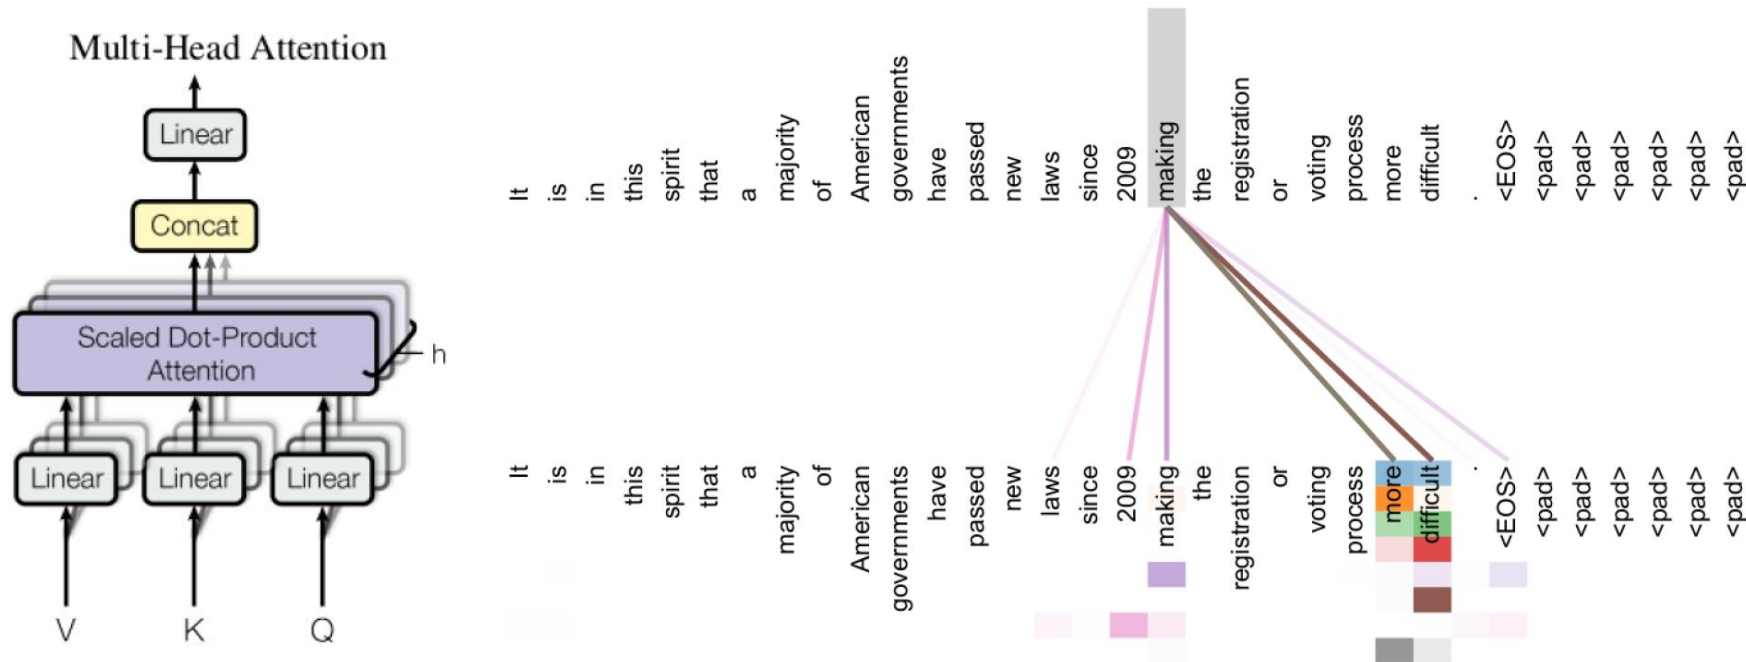
Vaswani et al. (2017) - Attention Is All You Need

- Input **embedding**: vocabulary-sized embedding
- **Positional encodings**: joins chunk index position
- Repeat N times
 - **Multi-headed attention** layer
 - Bypass & add input embedding, normalise
 - Two linear transformations & ReLU activation
 - Bypass & add embeddings, normalise
 - Feed forward to encoder
- Note: Model is bidirectional because next word predictions are done forwards and backwards



Paper

Vaswani et al. (2017) - Attention Is All You Need

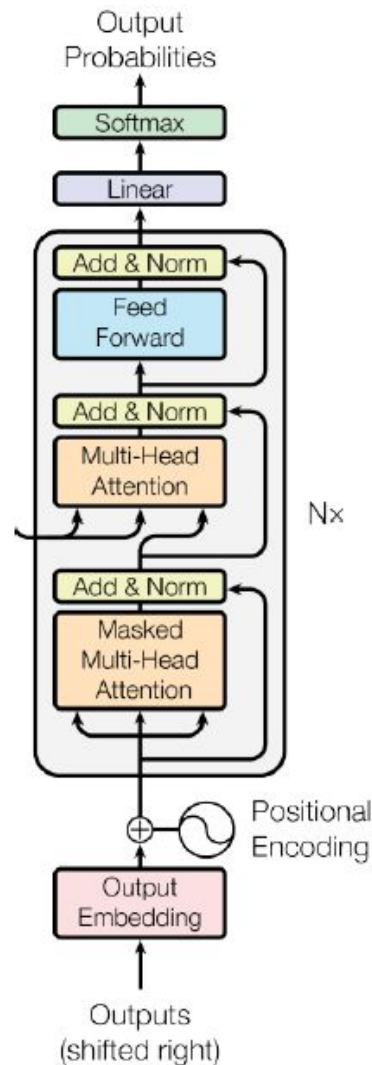


As nice as this is, this is also the models' main drawback: token limits

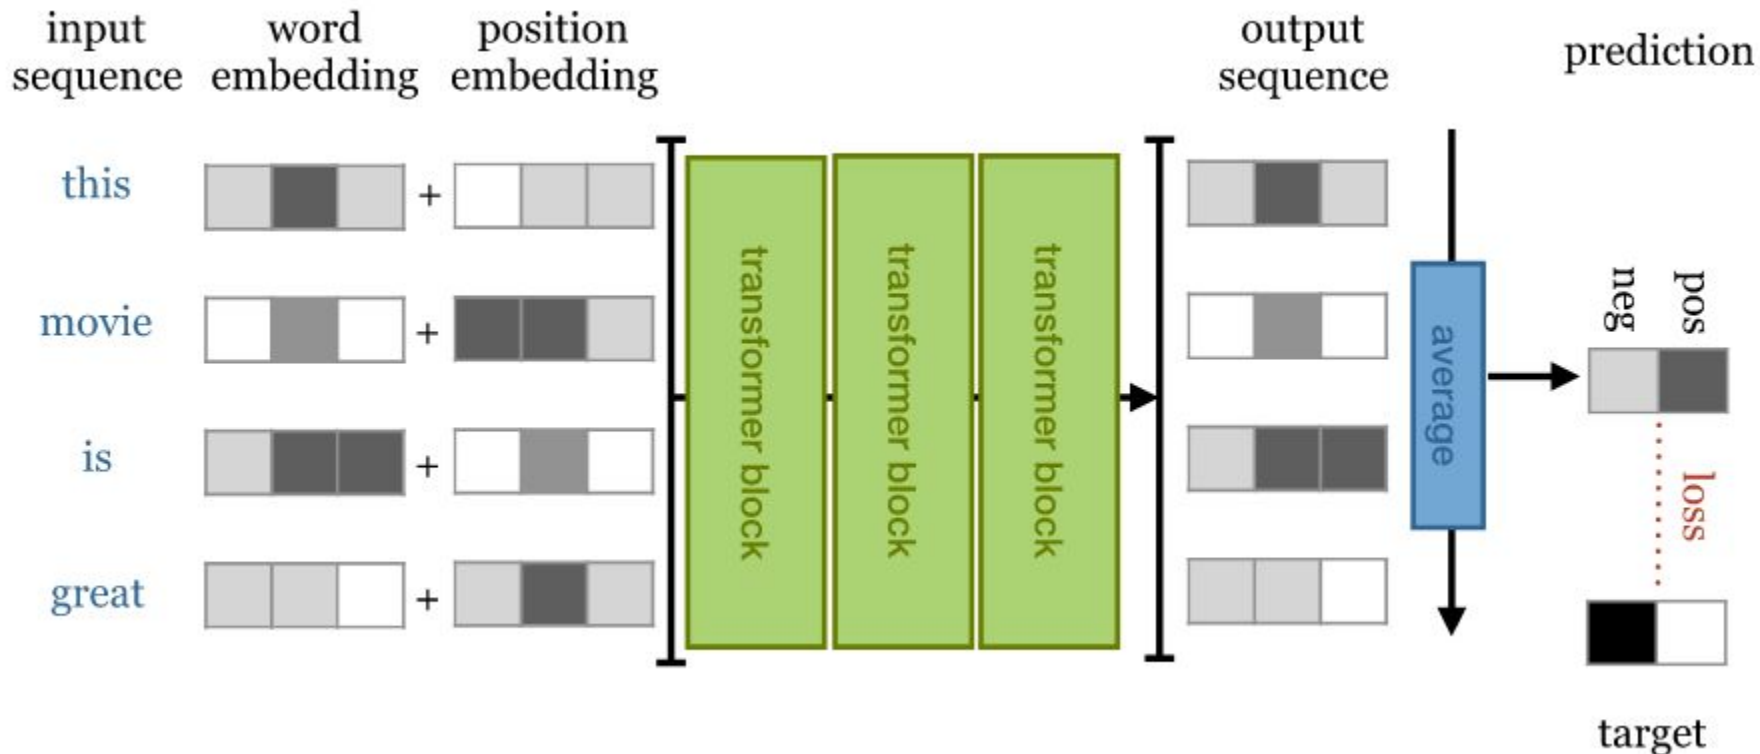
Paper

Vaswani et al. (2017) - *Attention Is All You Need*

- Output **embedding**: vocabulary-sized embedding
- **Positional encodings**: joins chunk index position
- Repeat N times
 - **Multi-headed attention** layer
 - Bypass & add input embedding, normalise
 - Multi-head attention for retrieving output embeddings using the encoder query
 - Bypass & add embeddings, normalise
 - Two linear transformations & ReLU activation
 - Bypass & add embeddings, normalise
- Linear transformation, softmax for probabilities
- Prediction



Transformers for practitioners: sentiment



Paper application

Blingler et al. (2023) - Cheap talk and cherry picking: ClimateBERT

Goal

Fine tune a BERT architecture to identify firms' voluntary disclosures related to climate (TCFD) & compare actions

Methodology

- **Create sample:** collect firms' annual reports and identify manually sentences related to climate disclosures align in one of several TCFD categories
- **Fine-tune model:** use a BERT decoder, and add Linear Layers on top to predict sentences categories (if apply)
- **Predict & Compare:** Compare talk with actual actions

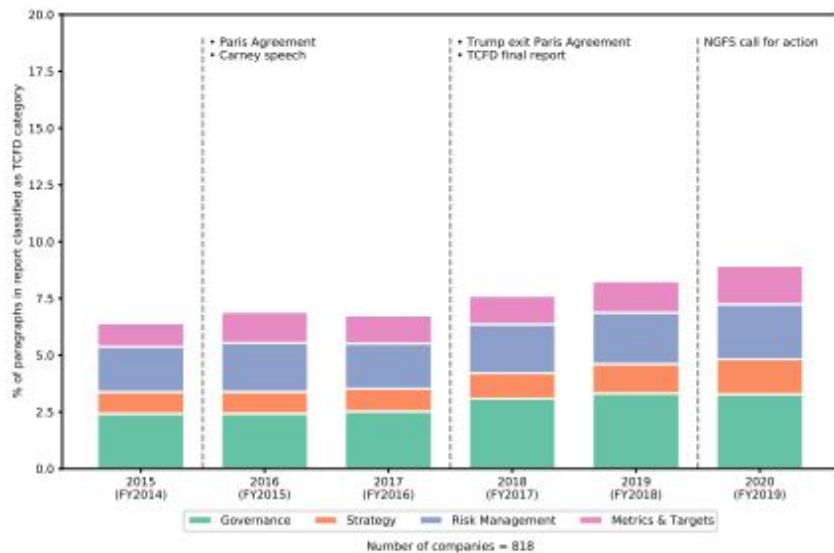
Results

- Voluntary disclosures suffer from **cheap talk**
- Announced TCFD support does not lead to an increase in disclosures, and cherry pick disclosures to overwhelmingly reflect non-material categories

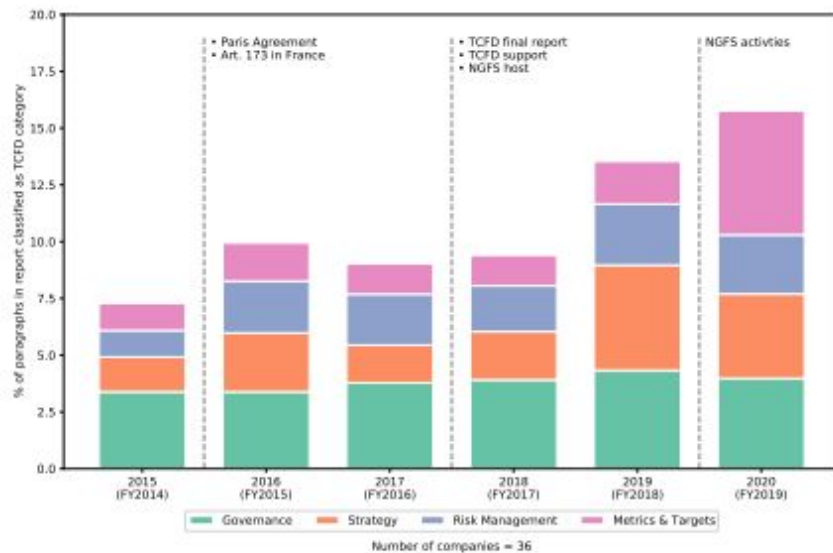
Governance	Strategy	Risk Management	Metrics and Targets
Disclose the organization's governance around climate-related risks and opportunities.	Disclose the actual and potential impacts of climate-related risks and opportunities on the organization's businesses, strategy, and financial planning where such information is material.	Disclose how the organization identifies, assesses, and manages climate-related risks.	Disclose the metrics and targets used to assess and manage relevant climate-related risks and opportunities where such information is material.
Recommended Disclosures			
Describe the board's oversight of climate-related risks and opportunities	Describe the climate-related risks and opportunities the organization has identified over the short, medium, and long term.	Describe the organization's processes for identifying and assessing climate-related risks	Disclose the metrics used by the organization to assess climate-related risks and opportunities in line with its strategy and risk management process.
Describe management's role in assessing and managing climate-related risks and opportunities.	Describe the impact of climate-related risks and opportunities on the organization's businesses, strategy, and financial planning.	Describe the organization's processes for managing climate-related risks.	Disclose Scope 1, Scope 2, and, if appropriate, Scope 3 greenhouse gas (GHG) emissions, and the related risks.
	Describe the resilience of the organization's strategy, taking into consideration different climate-related scenarios, including a 2°C or lower scenario.	Describe how processes for identifying, assessing, and managing climate-related risks are integrated into the organization's overall risk management.	Describe the targets used by the organization to manage climate-related risks and opportunities and performance against targets.

Paper application

Blingler et al. (2023) - Cheap talk and cherry picking: ClimateBERT

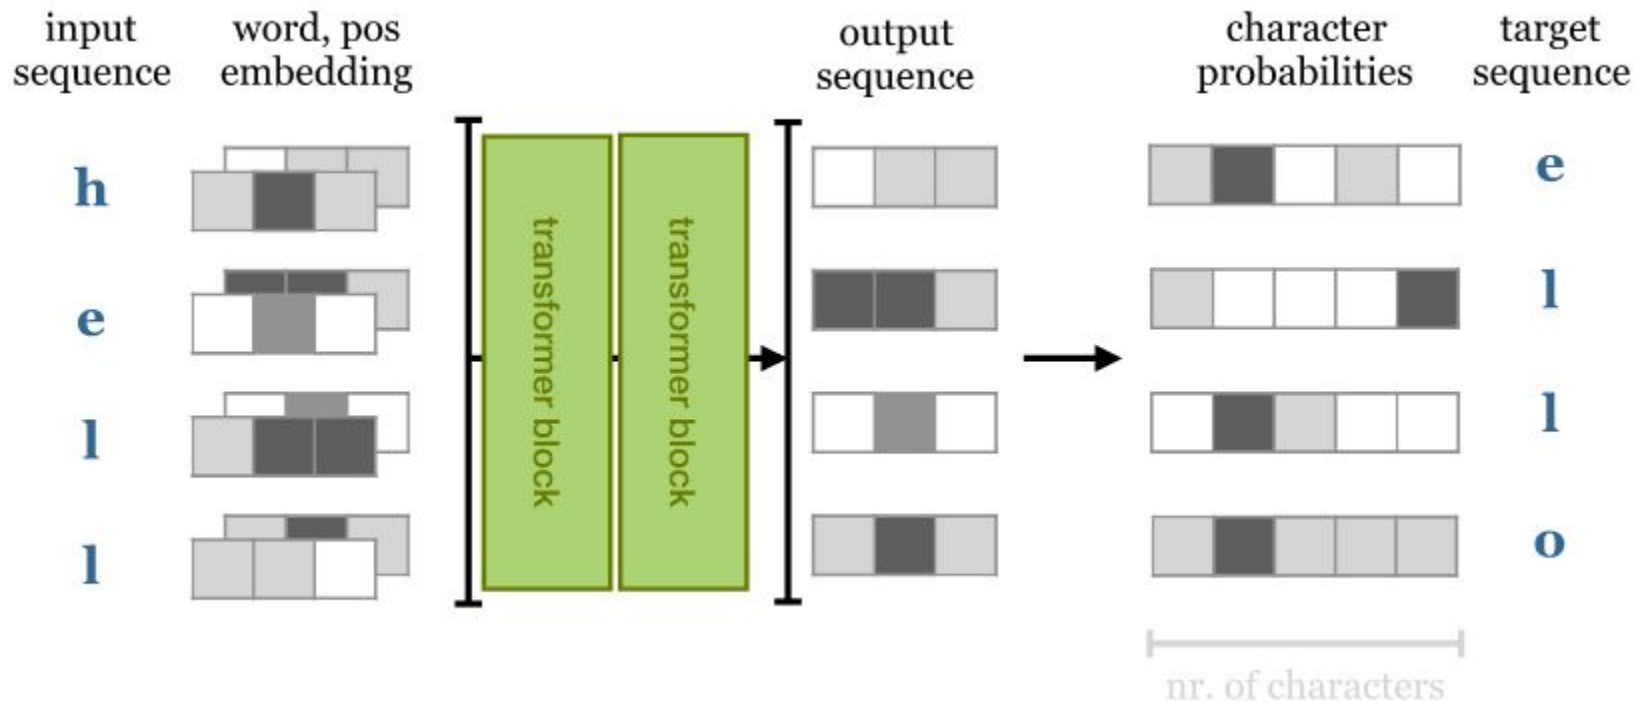


(a) Full sample



(b) France

Transformers for practitioners: autoregression



Evolution of Generative Pre-Trained Transformers

GPT-1 (2018) - 117M params

- Trained on the Books corpus
- Trained on a language modeling task, as well as a multi-task that adds a supervised learning task

GPT-2 (2019) - 1.5B params

- All articles linked from Reddit with at least 3 upvotes (8 million documents, 40GB of text)
- Dispense with supervised learning task, make adjustments, enlarge the model by many orders of magnitude

GPT-3 (2020) - 175B params

- Use an even larger corpus (Common Crawl, WebText2, Books1, Books2, Wikipedia)
- Model becomes much larger

InstructGPT ↔ GPT-3.5 ↔ GPT-4

- Add reinforcement learning with human feedback to improve responsiveness & user satisfaction
- Closed-source, controversial for OpenAI - an organisation built on the backs of the open-source crowd